

# M2 - Optimization and Inference Techniques for CV

## 2. Poisson editing

• • •

Group 8

Alex Carrillo  
Johnny Nuñez  
Guillem Martínez

MSc in Computer Vision

# Recap: Week 1

## Inpainting:

- Technique in which damaged, deteriorated or missing regions of images are reconstructed by interpolation of surrounding areas



## Conclusions and Results

- Notable results for inpainting small regions
- Bigger region, bigger blur in inpainted region, (may occur because):
  - The content available for inpainting is partial
  - The # of neighbours pixels to look at is not enough (4-connectivity)
- The smooth transition of the Laplacian operator is not enough to “hide” a visible region in an optimal way



# Poisson editing: this week's goal

Image A (*destination*)



“copy & paste”



Image B (*source*)



Insert one image (B) into another (A)...    ...without introducing any visually unappealing seams

OPTIMAL



# Criteria for the blending (Method #1: Importing Gradients)

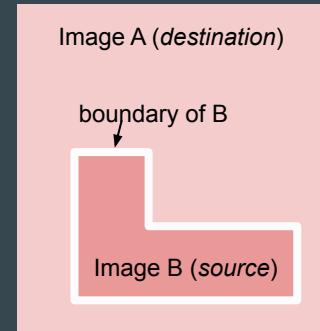
For a seamless “copy & paste” of the **inserted image (source; B)** into the **target image (destination; A)**, we want to generate a **new image H** to be pasted, that (visually):

1. Ensures that the color of the inserted image is shifted as part of the environment of the **target image**
2. Keeps intact all of the details (edges, corners, transitions, etc.) of the **inserted image**

## Preliminary intuition

This (mathematically thinking) translates into:

1. Pixels on H’s boundary should be the same as the pixels of A on that boundary
2. The gradient on the interior of H should be equal to the gradient on the interior of B



# Mathematically (Method #1: Importing Gradients)

For a seamless “copy & paste” of the **inserted image** (*source*; B) into the **target image** (*destination*; A), we want to generate a **new image H** to be pasted, that (visually):

1.  $H(x, y) = A(x, y)$  at each  $(x, y)$  of the boundary of B

2.  $\nabla H(x, y) = \nabla B(x, y)$  at each  $(x, y)$  inside B

$\Rightarrow$

$$4H(x, y) - \sum H(x + dx, y + dy) = 4B(x, y) - \sum B(x + dx, y + dy) \text{ at each } (x, y) \text{ inside B}$$

# Finite differences and Laplace operator (gradient computation)

Approximate gradients by means of finite differences equations, on each direction ( $x, y$ ).

Partial derivative	$\frac{\partial f(x)}{\partial x} \approx \Delta f(x) = \frac{f(x+h) - f(x)}{h} = [h := 1] = f(x+1) - f(x)$	Forward difference	Implemented in functions <i>DiFwd()</i> and <i>DjFwd()</i>
Partial derivative	$\frac{\partial f(x)}{\partial x} \approx \Delta f(x) = \frac{f(x) - f(x-h)}{h} = [h := 1] = f(x) - f(x-1)$	Backward difference	Implemented in functions <i>DiBwd()</i> and <i>DjBwd()</i>
2nd partial derivative	$\frac{\partial^2 f(x)}{\partial x^2} \approx \frac{\Delta^2 f(x)}{h^2} = \frac{\frac{f(x+h) - f(x)}{h} - \frac{f(x) - f(x-h)}{h}}{h} = \underbrace{f(x+1) - f(x)}_{\text{Forward difference}} - \underbrace{f(x) - f(x-1)}_{\text{Backward difference}}$		
	In the code, we can use: <i>drivingGrad_i</i> = DiBwd( $\cdot$ ) – DiFwd( $\cdot$ ) <i>drivingGrad_j</i> = DjBwd( $\cdot$ ) – DjFwd( $\cdot$ )	Forward difference Backward difference	
Laplacian operator	$\Delta f(x) = \sum_{i=1}^n \partial^2 f(x) = [\text{in 2D}] = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$	<pre>drivingGrad_i = G1_DiBwd( src(:,:,nC), param.hi ) - G1_DiFwd( src(:,:,nC), param.hi ); drivingGrad_j = G1_DjBwd( src(:,:,nC), param.hj ) - G1_DjFwd( src(:,:,nC), param.hj ); driving_on_src = drivingGrad_i + drivingGrad_j;</pre>	
	In the code, we can use: <i>driving_on_src/dst</i> = <i>drivingGRad_i</i> + <i>drivingGrad_j</i>		6

# Code implementation (Method #1: Importing Gradients)

1. Calculate the partial derivatives on each direction for the source image and calculate the central difference

```
% Second order central difference (second parcial derivatives)
drivingGrad_i = sol_DiBwd( src(:,:,nC), param.hi ) - sol_DiFwd( src(:,:,nC), param.hi );
drivingGrad_j = sol_DjBwd( src(:,:,nC), param.hj ) - sol_DjFwd( src(:,:,nC), param.hj );
```

2. Assemble the Laplacian operator by combining the central differences in each axis

```
% Laplacian
driving_on_src = drivingGrad_i + drivingGrad_j;
```

3. Create the gradients mask and compute the Poisson equation

```
driving_on_dst = zeros(size(dst(:,:,1)));
driving_on_dst(mask_dst(:)) = driving_on_src(mask_src(:));

param.driving = driving_on_dst;

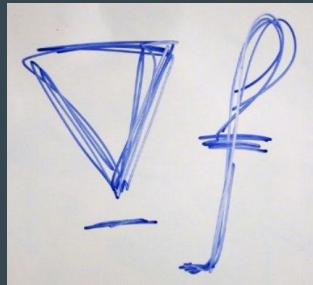
dst1(:,:,:,nC) = sol_Poisson_Equation_Axb(dst(:,:,:,nC), mask_dst, param);
```

# Goal (Method #2: Mixing Gradients)

With Importing Gradients, no trace of the destination image A is kept inside the source image B.

There are situations where it is desirable to **combine properties of A with those of B**

(i.e. add objects with holes, or partially transparent ones, on top of a textured or cluttered background).



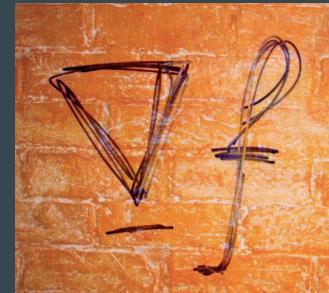
Source image B



Destination image A



Importing Gradients



Mixing Gradients

# Criteria for the blending (Method #2: Mixing Gradients)

For a seamless “copy & paste” of the **inserted image (source; B)** into the **target image (destination; A)**, we want to generate a **new image H** to be pasted, that (visually):

1. Ensures that the color of the inserted image is shifted as part of the environment of the **target image**  
~~Keeps intact all of the details (edges, corners, transitions, etc.) of the **inserted image**~~
2. Keeps the most relevant (striking) details of the each, the **inserted image** or the **target image**

## Preliminary intuition

This (mathematically thinking) translates into:

1. Pixels on H’s boundary should be the same as the pixels of A on that boundary
2. The interior of H should be equal to the maximum gradient (magnitude) between A and B

# Mathematically (Method #2: Mixing Gradients)

For a seamless “copy & paste” of the **inserted image** (*source*; B) into the **target image** (*destination*; A), we want to generate a **new image H** to be pasted, that (visually):

1.  $H(x, y) = A(x, y)$  at each  $(x, y)$  of the boundary of B
2.  $\nabla H(x, y) = \nabla A(x, y)$  if  $|\nabla A(x, y)| > |\nabla B(x, y)|$ , otherwise  $\nabla H(x, y) = \nabla B(x, y)$   
at each  $(x, y)$  inside B

# Code implementation (Method #2: Mixing Gradients)

1. Calculate the partial derivatives on each direction for source and destination

```
% Calculate the finite partial derivatives of the src
DiBwd_src = sol_DiBwd( src(:,:,nC), param.hi );
DiFwd_src = sol_DiFwd( src(:,:,nC), param.hi );
DjBwd_src = sol_DjBwd( src(:,:,nC), param.hj );
DjFwd_src = sol_DjFwd( src(:,:,nC), param.hj );\n\n% Calculate the finite partial derivatives of the dst
DiBwd_dst = sol_DiBwd( dst(:,:,nC), param.hi );
DiFwd_dst = sol_DiFwd( dst(:,:,nC), param.hi );
DjBwd_dst = sol_DjBwd( dst(:,:,nC), param.hj );
DjFwd_dst = sol_DjFwd( dst(:,:,nC), param.hj );
```

3. Calculate the central differences and assemble both directions to form the Laplacian operator

```
drivingGrad_i = DiBwd_dst-DiFwd_dst;
drivingGrad_j = DjBwd_dst-DjFwd_dst;
driving_on_dst = drivingGrad_i + drivingGrad_j;
```

2. Get the maximum gradient for each direction

```
% Get the maximum gradient on each direction
DiBwd_dst = get_max_grad(DiBwd_dst, DiBwd_src, mask_dst, mask_src);
DiFwd_dst = get_max_grad(DiFwd_dst, DiFwd_src, mask_dst, mask_src);
DjBwd_dst = get_max_grad(DjBwd_dst, DjBwd_src, mask_dst, mask_src);
DjFwd_dst = get_max_grad(DjFwd_dst, DjFwd_src, mask_dst, mask_src);
```

```
function [ result ] = get_max_grad( grad_dst, grad_src, mask_dst, mask_src )
% Compute the Forward finite differences with respect to the
% i coordinate only for the 1:end-1 rows. The last row is not replace
```

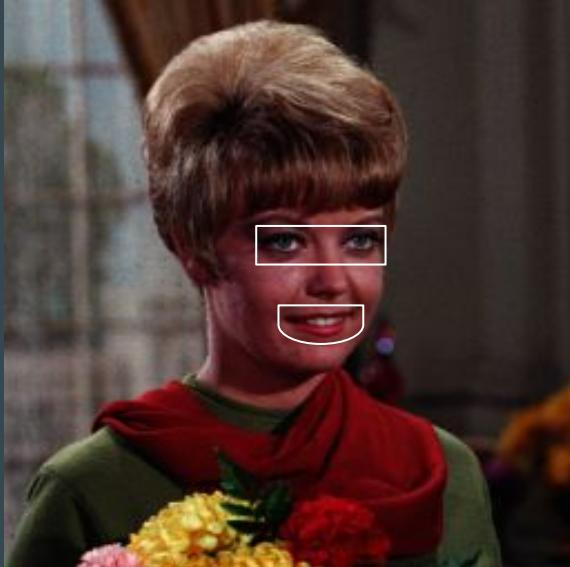
```
grad_dst_masked = grad_dst(mask_dst(:));
grad_src_masked = grad_src(mask_src(:));
grad_mask = abs(grad_dst_masked) <= abs(grad_src_masked);
grad_dst_masked(grad_mask) = grad_src_masked(grad_mask);
grad_dst(mask_dst(:)) = grad_dst_masked;
result = grad_dst;
end
```

4. Compute the Poisson equation

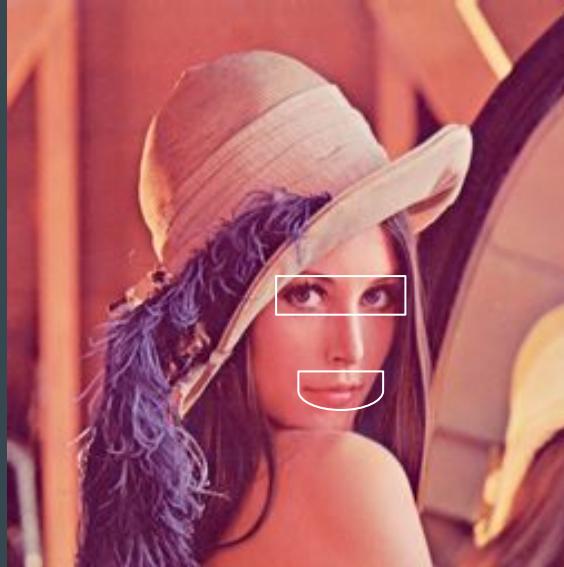
```
param.driving = driving_on_dst;
dst1(:,:,:,nC) = sol_Poisson_Equation_Axb(dst(:,:,:,nC), mask_dst, param);
```

# Results: Importing Gradients

Source



Target

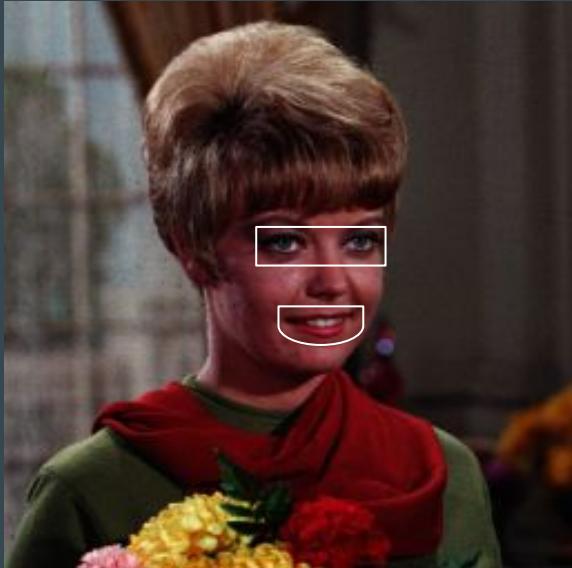


Result

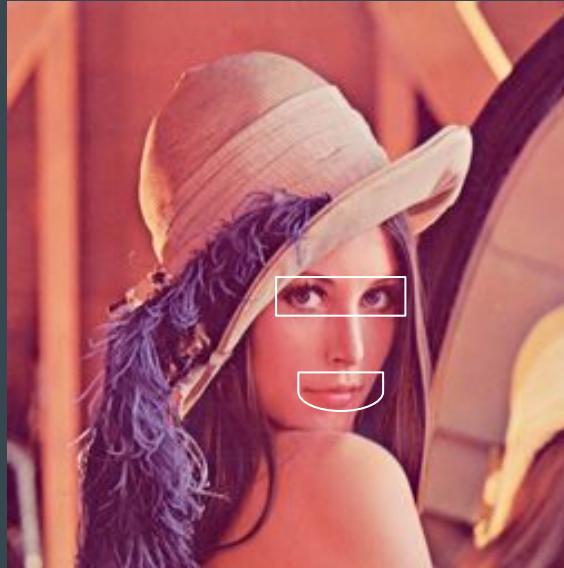


# Results: Mixing Gradients

Source



Target



Result



# Results: Importing Gradients

Source



Target



Result



# Results: Mixing Gradients

Source



Target



Result



# Custom Results: Importing Gradients

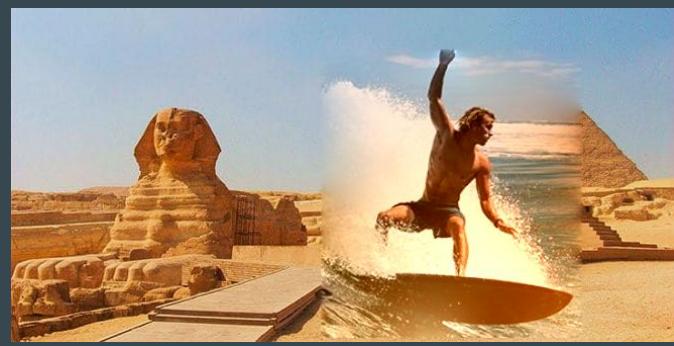
Source



Target



Result



# Custom Results: Mixing Gradients

Source



Target

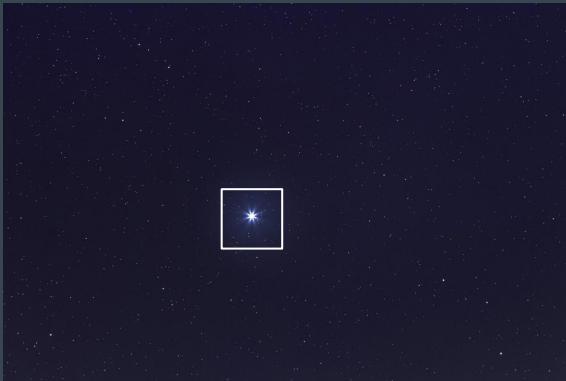


Result



# Custom Results: Importing Gradients

Source



Target



Result



# Custom Results: Mixing Gradients

Source



Target



Result



# Custom Results: Importing Gradients

Source



Target



Result



# Custom Results: Mixing Gradients

Source



Target



Result



# Custom Results: Importing Gradients

Source



Target



Result



# Custom Results: Mixing Gradients

Source



Target



Result



# Conclusions

- When pasting objects, importing gradients can be a better approach if the mask is defined properly delimiting the object.
- We obtain good results with Poisson, but there is a blurring in the background, on the cloned image.
- With Mixing gradients we can find unwanted textures when the pasted object is smooth and the destination image is not.
- The issue with the borders of the imported object is resolved by **mixed gradients**.
- When selecting the maximum, the system will not always work, because if the values you want to copy are smaller, you will lose information about the cloned image.